

# Import Data

```
from pyspark.sql import SparkSession
sc = SparkSession.builder.master("local[*]").getOrCreate()
```

```
data = spark.read.csv("/FileStore/tables/Bike_Rental_UCI_dataset-
bb6c6.csv",inferSchema=True,header=True)
```

```
data.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
----+-----+-----+
|season| yr|mnth|  hr|holiday|workingday|weathersit|temp|  hum|windspeed|dayOf
Week|days|demand|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
----+-----+-----+
|      1|  0|   1|   0|      0|      0|      1|0.24|0.81|      0.0|
Sat|   0|   16|
|      1|  0|   1|   1|      0|      0|      1|0.22| 0.8|      0.0|
Sat|   0|   40|
|      1|  0|   1|   2|      0|      0|      1|0.22| 0.8|      0.0|
Sat|   0|   32|
|      1|  0|   1|   3|      0|      0|      1|0.24|0.75|      0.0|
Sat|   0|   13|
|      1|  0|   1|   4|      0|      0|      1|0.24|0.75|      0.0|
Sat|   0|    1|
|      1|  0|   1|   5|      0|      0|      2|0.24|0.75| 0.0896|
Sat|   0|    1|
|      1|  0|   1|   6|      0|      0|      1|0.22| 0.8|      0.0|
Sat|   0|    2|
|      1|  0|   1|   7|      0|      0|      1| 0.2|0.86|      0.0|
Sat|   0|    3|
|      1|  0|   1|   8|      0|      0|      1|0.24|0.75|      0.0|
Sat|   0|    8|
|      1|  0|   1|   9|      0|      0|      1|0.32|0.76|      0.0|
Sat|   0|   14|
|      1|  0|   1|  10|      0|      0|      1|0.38|0.76| 0.2537|
Sat|   0|   36|
|      1|  0|   1|  11|      0|      0|      1|0.36|0.81| 0.2836|
Sat|   0|   56|
|      1|  0|   1|  12|      0|      0|      1|0.42|0.77| 0.2836|
Sat|   0|   84|
|      1|  0|   1|  13|      0|      0|      2|0.46|0.72| 0.2985|
```

```

Sat| 0| 94|
| 1| 0| 1| 14| 0| 0| 2|0.46|0.72| 0.2836|
Sat| 0| 106|
| 1| 0| 1| 15| 0| 0| 2|0.44|0.77| 0.2985|
Sat| 0| 110|
| 1| 0| 1| 16| 0| 0| 2|0.42|0.82| 0.2985|
Sat| 0| 93|
| 1| 0| 1| 17| 0| 0| 2|0.44|0.82| 0.2836|
Sat| 0| 67|
| 1| 0| 1| 18| 0| 0| 3|0.42|0.88| 0.2537|
Sat| 0| 35|
| 1| 0| 1| 19| 0| 0| 3|0.42|0.88| 0.2537|
Sat| 0| 37|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
only showing top 20 rows

```

# Feature Engineering

```

from pyspark.ml.classification import
(LogisticRegression,DecisionTreeClassifier,RandomForestClassifier)
from pyspark.ml import Pipeline
from pyspark.ml.tuning import ParamGridBuilder,CrossValidator
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.feature import StringIndexer

```

```

indexer = StringIndexer(inputCol="dayOfWeek",outputCol="day_cat")

```

```

indexed_data = indexer.fit(data).transform(data)
indexed_data.show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
|season| yr|mnth| hr|holiday|workingday|weathersit|temp| hum|windspeed|dayOf
Week|days|demand|day_cat|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
| 1| 0| 1| 0| 0| 0| 1|0.24|0.81| 0.0|
Sat| 0| 16| 0.0|
| 1| 0| 1| 1| 0| 0| 1|0.22| 0.8| 0.0|
Sat| 0| 40| 0.0|
| 1| 0| 1| 2| 0| 0| 1|0.22| 0.8| 0.0|
Sat| 0| 32| 0.0|
| 1| 0| 1| 3| 0| 0| 1|0.24|0.75| 0.0|
Sat| 0| 13| 0.0|

```



```
from pyspark.ml.feature import VectorAssembler
```

```
vec = VectorAssembler(inputCols=[
    'season','yr','mnth','hr','holiday','workingday','weathersit','temp','hum',
    'windspeed','days','day_cat'],outputCol='features')
```

```
data1 = vec.transform(indexed_data)
```

```
data1.show()
```

```
+-----+---+-----+---+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|season| yr|mnth| hr|holiday|workingday|weathersit|temp| hum|windspeed|dayOf
Week|days|demand|day_cat|          features|
+-----+---+-----+---+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      1| 0|   1| 0|      0|      0|      1|0.24|0.81|      0.0|
Sat|    0|  16|  0.0|(12,[0,2,6,7,8],[...|
|      1| 0|   1| 1|      0|      0|      1|0.22| 0.8|      0.0|
Sat|    0|  40|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 2|      0|      0|      1|0.22| 0.8|      0.0|
Sat|    0|  32|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 3|      0|      0|      1|0.24|0.75|      0.0|
Sat|    0|  13|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 4|      0|      0|      1|0.24|0.75|      0.0|
Sat|    0|   1|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 5|      0|      0|      2|0.24|0.75| 0.0896|
Sat|    0|   1|  0.0|[1.0,0.0,1.0,5.0,...|
|      1| 0|   1| 6|      0|      0|      1|0.22| 0.8|      0.0|
Sat|    0|   2|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 7|      0|      0|      1| 0.2|0.86|      0.0|
Sat|    0|   3|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 8|      0|      0|      1|0.24|0.75|      0.0|
Sat|    0|   8|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1| 9|      0|      0|      1|0.32|0.76|      0.0|
Sat|    0|  14|  0.0|(12,[0,2,3,6,7,8]...|
|      1| 0|   1|10|      0|      0|      1|0.38|0.76| 0.2537|
Sat|    0|  36|  0.0|[1.0,0.0,1.0,10.0...|
|      1| 0|   1|11|      0|      0|      1|0.36|0.81| 0.2836|
Sat|    0|  56|  0.0|[1.0,0.0,1.0,11.0...|
|      1| 0|   1|12|      0|      0|      1|0.42|0.77| 0.2836|
Sat|    0|  84|  0.0|[1.0,0.0,1.0,12.0...|
|      1| 0|   1|13|      0|      0|      2|0.46|0.72| 0.2985|
Sat|    0|  94|  0.0|[1.0,0.0,1.0,13.0...|
|      1| 0|   1|14|      0|      0|      2|0.46|0.72| 0.2836|
Sat|    0| 106|  0.0|[1.0,0.0,1.0,14.0...|
|      1| 0|   1|15|      0|      0|      2|0.44|0.77| 0.2985|
Sat|    0| 110|  0.0|[1.0,0.0,1.0,15.0...|
```

```

|      1|  0|   1| 16|      0|      0|      2|0.42|0.82|   0.2985|
Sat|    0|   93|   0.0|[1.0,0.0,1.0,16.0...|
|      1|  0|   1| 17|      0|      0|      2|0.44|0.82|   0.2836|
Sat|    0|   67|   0.0|[1.0,0.0,1.0,17.0...|
|      1|  0|   1| 18|      0|      0|      3|0.42|0.88|   0.2537|
Sat|    0|   35|   0.0|[1.0,0.0,1.0,18.0...|
|      1|  0|   1| 19|      0|      0|      3|0.42|0.88|   0.2537|
Sat|    0|   37|   0.0|[1.0,0.0,1.0,19.0...|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
only showing top 20 rows

```

```
modelData = data1.select('features','demand')
```

```
modelData.describe().show()
```

```

+-----+-----+
|summary|          demand|
+-----+-----+
|  count|          17379|
|   mean|189.46308763450142|
| stddev|181.3875990918646|
|    min|              1|
|    max|             977|
+-----+-----+

```

```
modelData.show(truncate=False)
```

```

+-----+-----+
|features|          demand|
+-----+-----+
|(12,[0,2,6,7,8],[1.0,1.0,1.0,0.24,0.81])|16|
|(12,[0,2,3,6,7,8],[1.0,1.0,1.0,1.0,0.22,0.8])|40|
|(12,[0,2,3,6,7,8],[1.0,1.0,2.0,1.0,0.22,0.8])|32|
|(12,[0,2,3,6,7,8],[1.0,1.0,3.0,1.0,0.24,0.75])|13|
|(12,[0,2,3,6,7,8],[1.0,1.0,4.0,1.0,0.24,0.75])|1|
|[1.0,0.0,1.0,5.0,0.0,0.0,2.0,0.24,0.75,0.0896,0.0,0.0]|1|
|(12,[0,2,3,6,7,8],[1.0,1.0,6.0,1.0,0.22,0.8])|2|
|(12,[0,2,3,6,7,8],[1.0,1.0,7.0,1.0,0.2,0.86])|3|
|(12,[0,2,3,6,7,8],[1.0,1.0,8.0,1.0,0.24,0.75])|8|
|(12,[0,2,3,6,7,8],[1.0,1.0,9.0,1.0,0.32,0.76])|14|
|[1.0,0.0,1.0,10.0,0.0,0.0,1.0,0.38,0.76,0.2537,0.0,0.0]|36|
|[1.0,0.0,1.0,11.0,0.0,0.0,1.0,0.36,0.81,0.2836,0.0,0.0]|56|
|[1.0,0.0,1.0,12.0,0.0,0.0,1.0,0.42,0.77,0.2836,0.0,0.0]|84|
|[1.0,0.0,1.0,13.0,0.0,0.0,2.0,0.46,0.72,0.2985,0.0,0.0]|94|
|[1.0,0.0,1.0,14.0,0.0,0.0,2.0,0.46,0.72,0.2836,0.0,0.0]|106|
|[1.0,0.0,1.0,15.0,0.0,0.0,2.0,0.44,0.77,0.2985,0.0,0.0]|110|

```

```
| [1.0,0.0,1.0,16.0,0.0,0.0,2.0,0.42,0.82,0.2985,0.0,0.0] |93 |
| [1.0,0.0,1.0,17.0,0.0,0.0,2.0,0.44,0.82,0.2836,0.0,0.0] |67 |
| [1.0,0.0,1.0,18.0,0.0,0.0,3.0,0.42,0.88,0.2537,0.0,0.0] |35 |
| [1.0,0.0,1.0,19.0,0.0,0.0,3.0,0.42,0.88,0.2537,0.0,0.0] |37 |
+-----+-----+
only showing top 20 rows
```

## Seperate data for Train & Test

```
trainData,testData = modelData.randomSplit([0.7,0.3])
```

```
trainData = trainData.withColumnRenamed(('demand'),('label'))
```

```
testData = testData.withColumnRenamed(('demand'), ('label'))
```

```
trainData.show(truncate=False)
```

```
+-----+-----+
| features                                | label |
+-----+-----+
| (12,[0,2,3,6,7,8],[1.0,1.0,1.0,1.0,0.22,0.8]) |40 |
| (12,[0,2,3,6,7,8],[1.0,1.0,2.0,1.0,0.22,0.8]) |32 |
| (12,[0,2,3,6,7,8],[1.0,1.0,3.0,1.0,0.24,0.75]) |13 |
| (12,[0,2,3,6,7,8],[1.0,1.0,6.0,1.0,0.22,0.8]) |2 |
| (12,[0,2,3,6,7,8],[1.0,1.0,8.0,1.0,0.24,0.75]) |8 |
| (12,[0,2,3,6,7,8],[1.0,1.0,9.0,1.0,0.32,0.76]) |14 |
| (12,[0,2,6,7,8,10],[1.0,1.0,1.0,0.18,0.55,13.0]) |28 |
| (12,[0,2,6,7,8,10],[2.0,4.0,2.0,0.3,0.61,87.0]) |32 |
| (12,[0,2,6,7,8,10],[2.0,6.0,1.0,0.56,0.52,150.0]) |93 |
| (12,[0,2,6,7,8,10],[3.0,7.0,1.0,0.82,0.56,199.0]) |101 |
| (12,[0,2,6,7,8,10],[3.0,8.0,1.0,0.6,0.88,227.0]) |128 |
| (12,[0,2,6,7,8,10],[4.0,10.0,1.0,0.48,0.72,275.0]) |89 |
| [1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.04,0.45,0.2537,19.0,0.0] |13 |
| [1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.04,0.57,0.1045,20.0,1.0] |22 |
| [1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.22,0.64,0.3582,25.0,0.0] |28 |
| [1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.26,0.56,0.0,14.0,1.0] |39 |
| [1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.06,0.41,0.19399999999999998,21.0,3.0] |7 |
| [1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.12,0.5,0.19399999999999998,12.0,2.0] |14 |
| [1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.12,0.5,0.2836,8.0,3.0] |5 |
| [1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.14,0.59,0.2836,11.0,5.0] |7 |
+-----+-----+
```

```
only showing top 20 rows
```

```
testData.show(truncate=False)
```

```
+-----+-----+
| features                                | label |
+-----+-----+
|(12,[0,2,3,6,7,8],[1.0,1.0,4.0,1.0,0.24,0.75])| 1     |
|(12,[0,2,3,6,7,8],[1.0,1.0,7.0,1.0,0.2,0.86])| 3     |
|(12,[0,2,6,7,8],[1.0,1.0,1.0,0.24,0.81])| 16    |
|(12,[0,2,6,7,8,10],[3.0,9.0,1.0,0.62,0.94,247.0])| 116   |
|[1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.1,0.42,0.3881,7.0,1.0]| 25    |
|[1.0,0.0,1.0,0.0,0.0,0.0,1.0,0.16,0.8,0.1045,26.0,1.0]| 33    |
|[1.0,0.0,1.0,0.0,0.0,0.0,2.0,0.18,0.51,0.1642,6.0,0.0]| 25    |
|[1.0,0.0,1.0,0.0,0.0,0.0,2.0,0.46,0.88,0.2985,1.0,1.0]| 17    |
|[1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.14,0.59,0.1045,9.0,6.0]| 12    |
|[1.0,0.0,1.0,0.0,0.0,1.0,1.0,0.26,0.56,0.3881,17.0,5.0]| 13    |
|[1.0,0.0,1.0,0.0,0.0,1.0,2.0,0.2,0.75,0.1343,24.0,2.0]| 9     |
|[1.0,0.0,1.0,0.0,0.0,1.0,2.0,0.22,0.93,0.0,17.0,4.0]| 3     |
|[1.0,0.0,1.0,1.0,0.0,0.0,1.0,0.14,0.8,0.0,26.0,1.0]| 29    |
|[1.0,0.0,1.0,1.0,0.0,0.0,2.0,0.16,0.59,0.0896,13.0,0.0]| 20    |
|[1.0,0.0,1.0,1.0,0.0,1.0,2.0,0.2,0.69,0.2239,5.0,2.0]| 7     |
|[1.0,0.0,1.0,1.0,0.0,1.0,2.0,0.24,0.65,0.1343,23.0,4.0]| 5     |
|[1.0,0.0,1.0,1.0,0.0,1.0,3.0,0.22,0.93,0.1343,17.0,4.0]| 7     |
|[1.0,0.0,1.0,2.0,0.0,0.0,2.0,0.18,0.55,0.0,6.0,0.0]| 16    |
|[1.0,0.0,1.0,2.0,0.0,1.0,1.0,0.14,0.86,0.1343,10.0,4.0]| 1     |
|[1.0,0.0,1.0,2.0,0.0,1.0,1.0,0.16,0.64,0.0,4.0,5.0]| 2     |
+-----+-----+
```

only showing top 20 rows

# Linear Regression

```
from pyspark.ml.regression import LinearRegression
lr = LinearRegression(featuresCol='features', labelCol='label', maxIter=10,
regParam=0.3, elasticNetParam=1)
lr_model = lr.fit(trainData)
print("Coefficients: " + str(lr_model.coefficients))
print("Intercept: " + str(lr_model.intercept))
trainingSummary = lr_model.summary
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("r2: %f" % trainingSummary.r2)
```

```
Coefficients: [19.6914190944,53.0943621795,-2.09734576019,7.56361284259,-18.
6125249833,4.89479402101,-3.94093165934,282.052218461,-201.599320208,30.2210
596181,0.0748403062638,-0.265478966667]
Intercept: -2.5016762683676923
RMSE: 142.436392
r2: 0.385282
```

```
# Make predictions.
predictions = lr_model.transform(testData)

# Select example rows to display.
predictions.select("prediction", "label", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator1 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator1.evaluate(predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
evaluator2 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="r2")
r2 = evaluator2.evaluate(predictions)
print("R2 on test data = %g" % r2)
```

```
+-----+-----+-----+
|      prediction|label|      features|
+-----+-----+-----+
|-42.101040948732376|    1|(12,[0,2,3,6,7,8]...|
| -52.8682163823581|    3|(12,[0,2,3,6,7,8]...|
| -84.45145153157792|   16|(12,[0,2,6,7,8],[...|
|  37.61010761104589|  116|(12,[0,2,6,7,8,10...|
| -33.32783081995555|   25|[1.0,0.0,1.0,0.0,...|
+-----+-----+-----+
```

only showing top 5 rows

```
Root Mean Squared Error (RMSE) on test data = 140.911
R2 on test data = 0.39197
```

## LR Cross Validation

```
from pyspark.ml import Pipeline
from pyspark.ml.tuning import ParamGridBuilder,CrossValidator
from pyspark.ml.evaluation import RegressionEvaluator
```

```
evaluator = RegressionEvaluator(metricName = 'r2')
```

```
evaluator.explainParam('metricName')
```

```
Out[118]: 'metricName: metric name in evaluation - one of:\n
rmse - root mean squared error (default)\n                               mse - mean
squared error\n                               r2 - r^2 metric\n
mae - mean absolute error. (default: rmse, current: r2)'
```



```
pipeline = Pipeline(stages=[lr])
```

```
gridBuilder=ParamGridBuilder()\n    .addGrid(lr.regParam, [0.1, 0.01]) \n    .build()
```

```
cv =\nCrossValidator(estimator=pipeline,estimatorParamMaps=gridBuilder,evaluator=e\nvaluator,numFolds=2)
```

```
cvm = cv.fit(trainData)
```

```
predictions = cvm.transform(testData)
```

```
evaluator.evaluate(predictions)
```

```
Out[124]: 0.10068787153472136
```

The result is unexpected, we need more time to investigate the reason.

## Display with group by

```
data1.groupby('season').mean('demand').show()
```

```
+-----+-----+
| season | avg(demand) |
+-----+-----+
|      1 | 111.11456859971712 |
|      3 | 236.01623665480426 |
|      4 | 198.86885633270322 |
|      2 | 208.34406894987526 |
+-----+-----+
```

```
data1.groupby('hr').mean('demand').show()
```

```
+---+-----+
| hr | avg(demand) |
+---+-----+
| 12 | 253.31593406593407 |
| 22 | 131.33516483516485 |
```

```
| 1| 33.3756906077348|
| 13|253.66117969821673|
| 6| 76.04413793103448|
| 16| 311.9835616438356|
| 3|11.727403156384504|
| 20|226.03021978021977|
| 5| 19.88981868898187|
| 19|311.52335164835165|
| 15| 251.2331961591221|
| 9|219.30949105914718|
| 17|461.45205479452056|
| 4| 6.352941176470588|
| 8|359.01100412654745|
| 23| 87.83104395604396|
| 7| 212.0646492434663|
| 10| 173.6685006877579|
| 21|172.31456043956044|
| 11| 208.1430536451169|
```

```
+----+-----+
```

only showing top 20 rows

## Other Models

```

from pyspark.ml import Pipeline
from pyspark.ml.regression import DecisionTreeRegressor
from pyspark.ml.feature import VectorIndexer
from pyspark.ml.evaluation import RegressionEvaluator

# Train a DecisionTree model.
dt = DecisionTreeRegressor(featuresCol="features")

# Chain indexer and tree in a Pipeline
pipeline = Pipeline(stages=[dt])

# Train model. This also runs the indexer.
model = pipeline.fit(trainData)

# Make predictions.
predictions = model.transform(testData)

# Select example rows to display.
predictions.select("prediction", "label", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator1 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator1.evaluate(predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
evaluator2 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="r2")
r2 = evaluator2.evaluate(predictions)
print("R2 on test data = %g" % r2)

```

```

+-----+-----+-----+
| prediction|label| features|
+-----+-----+-----+
| 14.869767441860466| 1|(12,[0,2,3,6,7,8]...|
| 107.7| 3|(12,[0,2,3,6,7,8]...|
| 66.51470588235294| 16|(12,[0,2,6,7,8],[...|
| 124.61904761904762| 116|(12,[0,2,6,7,8,10...|
| 66.51470588235294| 25|[1.0,0.0,1.0,0.0,...|
+-----+-----+-----+

```

only showing top 5 rows

Root Mean Squared Error (RMSE) on test data = 107.384  
R2 on test data = 0.646884

```

from pyspark.ml import Pipeline
from pyspark.ml.regression import RandomForestRegressor
from pyspark.ml.feature import VectorIndexer
from pyspark.ml.evaluation import RegressionEvaluator

# Train a RandomForest model.
rf = RandomForestRegressor(featuresCol='features')

# Chain indexer and forest in a Pipeline
pipeline = Pipeline(stages=[rf])

# Train model. This also runs the indexer.
model = pipeline.fit(trainData)

# Make predictions.
predictions = model.transform(testData)

# Select example rows to display.
predictions.select("prediction", "label", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator1 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator1.evaluate(predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
evaluator2 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="r2")
r2 = evaluator2.evaluate(predictions)
print("R2 on test data = %g" % r2)

```

```

+-----+-----+-----+
|      prediction|label|      features|
+-----+-----+-----+
| 19.70270694650672|   1|(12,[0,2,3,6,7,8]...|
| 74.45332008243942|   3|(12,[0,2,3,6,7,8]...|
|30.302337702024925|  16|(12,[0,2,6,7,8],[...|
| 67.10513669994847| 116|(12,[0,2,6,7,8,10...|
|33.211464342172775|  25|[1.0,0.0,1.0,0.0,...|
+-----+-----+-----+

```

only showing top 5 rows

```

Root Mean Squared Error (RMSE) on test data = 110.492
R2 on test data = 0.626147

```

```

from pyspark.ml import Pipeline
from pyspark.ml.regression import GBRegressor
from pyspark.ml.feature import VectorIndexer
from pyspark.ml.evaluation import RegressionEvaluator

# Train a GBT model.
gbt = GBRegressor(featuresCol='features', maxIter=50)

# Chain indexer and GBT in a Pipeline
pipeline = Pipeline(stages=[gbt])

# Train model. This also runs the indexer.
model = pipeline.fit(trainData)

# Make predictions.
predictions = model.transform(testData)

# Select example rows to display.
predictions.select("prediction", "label", "features").show(5)

# Select (prediction, true label) and compute test error
evaluator1 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator1.evaluate(predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
evaluator2 = RegressionEvaluator(
    labelCol="label", predictionCol="prediction", metricName="r2")
r2 = evaluator2.evaluate(predictions)
print("R2 on test data = %g" % r2)

```

```

+-----+-----+-----+
|      prediction|label|      features|
+-----+-----+-----+
|-13.52827471851879|    1|(12,[0,2,3,6,7,8]...|
|-44.0988006791547|    3|(12,[0,2,3,6,7,8]...|
| 28.40194416751362|   16|(12,[0,2,6,7,8],[...|
|106.60344271916942|  116|(12,[0,2,6,7,8,10...|
|29.900054705811844|   25|[1.0,0.0,1.0,0.0,...|
+-----+-----+-----+

```

only showing top 5 rows

Root Mean Squared Error (RMSE) on test data = 48.446

R2 on test data = 0.928129

For this Bike Rental Dataset, the best model is gradient-boosted tree regression, the RMSE is just 48.446 and the r2 is up to 0.928129 !